

Diagnosing Memory-Write Decisions in Conversational Assistants Through Specification-Based Replay

ANONYMOUS AUTHOR(S)

Memory-write evaluations must distinguish completing requested changes from corrupting stored facts. We replay saved extractions and proposals under four admission rules, score the executed states, and count exposure: outputs on which the rules can decide differently before scoring. On 48 Korean cases and 21 model configurations (3,024 episodes), this analysis yields three findings. First, evidence checks recover 11 corrections on Gemini 2.5 Flash, but releasing all eligible held corrections recovers the same 11. The holds originate in empty extracted fact lists; a revised extraction instruction lets the original gate complete five of five selected corrections, against one of five under the original instruction. Second, 12 of 21 configurations have zero exposure, so their ties do not test decisions on which the rules differ. Third, accepting surface variants of required facts reduces the violations avoided by a gate on general-review proposals from 82 to 30 and increases lost completed changes from 40 to 92. Differences among the replayed correction rules remain unchanged under both references. On real speech-recognition hypotheses, the same trace instead locates holds in disagreement between hypotheses. These findings identify failure mechanisms and reference sensitivity in the tested pipeline; Section 8 discusses their scope and further validation.

CCS Concepts: • **General and reference** → **Evaluation**; • **Computing methodologies** → *Nonmonotonic, default reasoning and belief revision*; *Intelligent agents*; Discourse, dialogue and pragmatics.

Additional Key Words and Phrases: conversational agents, persistent memory, memory correction, admission control, specification-based evaluation, failure diagnosis


1 Introduction

Conversational assistants increasingly maintain persistent memories of their users [13]. Updating these memories requires completing requested changes while preserving facts that should remain unchanged. Systems propose additions, revisions, and deletions [4, 8, 29, 48, 49], but downstream answers or final-state scores [26, 46] do not identify which stage produced an error. A correction may be proposed correctly yet withheld because an earlier extraction omitted its supporting fact.

We study this problem by separating generation, admission, and execution. Generation produces extracted facts and proposed operations; admission determines which operations may execute; execution changes the stored state. Our specification-based replay holds the generated artifacts fixed, applies alternative admission rules, and checks the whole resulting state against an explicit reference. The four rules range from keeping corrections without shared extracted support to releasing eligible held corrections or requiring literal or target-linked evidence. Their *exposure* is the set of outputs on which they can decide differently, determined from gate logs before scoring.

The study uses 48 authored Korean cases, 21 model configurations, and three repetitions per case. Its three main findings are (Figure 1):

- *Extraction omissions can create apparent gains for admission rules.* On Gemini 2.5 Flash, all three rules that relax the agreement gate recover the same 11 corrections. Empty extracted fact lists created these holds. In a follow-up on five selected cases, revising the extraction instruction lets the unchanged gate complete five corrections, against one under the original instruction (Section 6.1).



Counts pool three repetitions. Reference violations follow the frozen state specification.

Fig. 1. Three evaluation confounds that a final-state score cannot distinguish: attribution (A), specification (B), and opportunity (C). Each row sets what the final-state score shows beside what replay of the same saved artifacts shows. A: Gemini 2.5 Flash, strict corrections completed out of 33 observations under each rule. B: The 146 reference-violating states under REVIEW on the 21-configuration panel, split by whether the gate avoids them and, if so, by the held operation; “variants” are append holds whose tuple differs from a required fact only in subject label, spacing, or value form (Section 6.4). C: The 21 panel configurations, with controlled-path and with REVIEW proposals; hollow squares mark configurations on which no proposal reached the held set H .

- *Many ties arise without a decision that could separate the rules.* Twelve of 21 configurations have zero exposure on controlled-path proposals. Across the panel, only 49 outputs from 12 cases reach the decision boundary. Exposure therefore determines which ties test competing admission decisions (Section 6.2).
- *Reference conventions change the apparent benefit of gating.* Accepting surface variants of required facts reduces avoided violations from 82 to 30 and increases lost completed changes from 40 to 92 on general-review proposals. Pairwise differences among the replayed correction rules remain unchanged under both references (Section 6.4).

These findings concern attribution, opportunity, and specification, respectively. For fixed generated artifacts and admission rules, exposure is determined from gate decisions before scoring, independently of the endpoint and reference, and bounds the rule differences that any scored comparison can reveal. Table 1 places this evaluation alongside eleven memory-write studies. The remaining analyses examine target-linked evidence, decompose the complete-pipeline comparison, and apply replay to a learned verifier and real recognition hypotheses (Sections 6.3 and 6.5–6.7). The findings establish mechanisms in the tested pipeline, not their prevalence across systems.

2 Related Work

2.1 Memory Updates and Admission in LLM Agents

Several systems give LLM agents a persistent memory that is updated through explicit operations. MemGPT treats the context window as fast memory and moves information to and from external storage through function calls [29]. Mem0 extracts candidate facts and applies one of four operations to existing memories: add, update, delete, or no operation [8]. A-Mem stores memories as linked notes whose attributes are revised as new notes arrive [48], and MemoryBank, recursive summarization, SeCom, LD-Agent, and Generative Agents maintain or select conversational history in related ways [20, 30, 31, 44, 58]. Closest to the correction decision studied here, Bae et al. cast memory management in Korean long-term conversation as pairwise operations between a stored sentence and a new session summary (pass, replace, append, or delete) [4]. They compared trained classifiers with natural language inference models mapped to these operations and evaluated operation accuracy, set-level overlap, and human ratings of later sessions. Memory-R1 learns add, update, delete, and no-operation decisions from a reward based on downstream answer correctness [49], and Zep invalidates, rather than deletes, knowledge-graph edges that newer information contradicts [34].

A recent line of work treats the write itself as an admission decision. A-MAC learns an admission policy from interpretable factors (future utility, factual confidence, semantic novelty, temporal recency, and a content-type prior) and evaluates admission decisions against labeled candidate memories [54]; ConsistencyGate samples the same model several times to score whether a candidate fact is supported by its source and admits the fact above a threshold [56]; and DerivAudit audits whether a stored memory follows from the history available when it was written, finding that support often lies beyond the writer’s citations and that verifiers still admit many unsupported memories [21]. STALE evaluates whether agents notice that a stored memory has been implicitly invalidated by later information [6]. Memory writes are also an attack surface: AgentPoison and MINJA show that poisoned or injected records, the latter inserted through ordinary queries, later steer agent behavior [7, 9]. In our setting, the surrounding history is supplied rather than retrieved, so retrieval failure is not confounded with admission failure; how the position of information affects its use is a separate question [22].

2.2 State Revision and Competing Hypotheses

Dialogue-state tracking already models state revision. SOM-DST predicts an operation for each slot before generating values [19], and its oracle analysis attributes most errors to the operation predictor rather than to value generation, an early stage-level diagnosis of the kind we apply to memory writes. AG-DST amends an initially generated state [42], dataflow synthesis supports executable dialogue programs and non-destructive revision [3], and PoA generates semantic-action code that updates the dialogue state [40]. Other trackers select the relevant turns, suppress incorrect values, or rewrite queries to resolve coreference [14, 23, 43]. The second Dialog State Tracking Challenge supplied competing speech-recognition hypotheses for every turn and scored the tracked state turn by turn [16]; its setting is close to ours, although there the tracker produces the state itself, and Section 6.7 replays our admission rules on its hypotheses. MultiWOZ 2.2 removes values that the system offered before the user accepted them [53], and MultiWOZ 2.4 removes values that only the system mentioned [51]; both revisions show why the timing and source of evidence matter. MultiWOZ 2.2 also lists the surface variants of a value as equally correct [53], so an exact-match reference can count an equivalent mention as an error (Section 6.4). Our comparison intervenes after extraction and proposal are fixed: it evaluates the permission to execute a proposal and thus complements these state-generation methods.

Table 1. Reporting audit of eleven memory-write evaluations, from what each paper reports. Included are the works of Section 2.1 that make or evaluate per-record decisions to add, revise, delete, forget, or admit memories and report a quantitative evaluation; systems that rewrite one summary or keep all history, and attack studies, are excluded. *Fixed*: the compared write or admission policies act on the same generated candidates. *Both errors*: missed required writes and wrongly admitted writes are reported, not only downstream answers. *Opportunity*: the number of cases on which the compared policies could, or did, differ; this article defines the quantity (exposure, Section 4.2), so a dash records only that it is not reported. *Matching*: how outputs are compared with references; n.r., not reported.

Work	Write decision scored	Fixed	Both errors	Opportunity	Matching
MemoryBank [58]	Downstream answers	–	–	–	Human ratings; agreement n.r.
MemGPT [29]	Downstream answers	–	–	–	LLM judge and ROUGE-L; agreement n.r.
Mem0 [8]	Downstream answers	–	–	–	Token overlap and LLM judge; agreement n.r.
A-Mem [48]	Downstream answers	–	–	–	Token and embedding overlap
Zep [34]	Downstream answers	–	–	–	LLM judge
Memory-R1 [49]	Downstream answers	–	–	–	Token overlap and LLM judge
STALE [6]	Downstream probes	–	Partial	–	LLM judge with human check
Bae et al. [4]	Operation per memory pair; set overlap	Yes	Partial	–	Human labels, majority of three; 87.5% annotator–gold agreement
A-MAC [54]	Admission per candidate	Yes	Yes	–	Candidate labels; process n.r.
ConsistencyGate [56]	Admission per fact	Partial	Yes	–	Constructed labels
DerivAudit [21]	Support and admission per memory	Yes	Yes	–	Two LLM annotators, blinded LLM adjudication; 96.0% human-review agreement
This work	Admission, whole state	Yes	Yes	Yes	Authored references, checked independently by two authors; exact and variant-tolerant matching

2.3 Evidence, Clarification, and Evaluation

Clarification has long been modeled as a decision under uncertainty [10, 12, 32, 52, 55], although human clarification is only weakly tied to annotated ambiguity [25] and is not mirrored by model uncertainty [41]; an uncertain detail can be unnecessary for the next action yet consequential for a durable write. Disagreeing readings resemble the inter-context conflicts studied between retrieved sources [47], and a quoted span shows that a value occurs in the source, not that it supports the proposed fact [1]; this distinction motivates comparing literal evidence with target-linked witnesses. Attribution to identified sources can be defined and rated by people [33], entailment-based checkers detect factual inconsistency [18], LLM judges can approximate human preferences [57], and agreement across readings resembles self-consistency over sampled outputs [45]. τ -bench scores agents by the final database state after tool use [50], the kind of executed-state endpoint that we apply to a single turn. Self-Refine and Reflexion revise outputs at inference time [24, 38]; we use a two-call general review as the alternative complete pipeline, and because the choice of instruction can itself change results [2], every condition receives the same task facts and retention constraints. Automatic and human assessments can agree [39] or rank systems differently [36], so agreement has to be established for each endpoint. Table 1 audits what memory-write evaluations report.

Only A-MAC and DerivAudit combine fixed candidates with both error types in this audit; none reports the opportunity for compared policies to differ. Our evaluation adds exposure and reference sensitivity to whole-state comparisons of edits to identified records under competing readings. We do not directly benchmark A-MAC, ConsistencyGate, or DerivAudit because their candidate-label tasks differ from the executed-state task here. Section 6.6 instead applies a learned verifier to the same saved proposals.

Table 2. Terms and endpoints. With three repetitions, the correction, positive, and all-case denominators are 33, 84, and 144 per configuration; with one repetition, they are 11, 28, and 48.

Term	Definition
Reading	One written alternative of the current utterance, supplied as text.
Extraction F	Reading-level fact lists and correction witnesses produced by the first model call.
Proposal (a, u, Q)	Action, reply, and memory operations produced by the second model call.
Held set H	CORRECT operations that AGREE holds only for lack of shared support, although the extraction validly covers every reading.
Exposure	Outputs on which the rules can decide differently, that is, outputs that contain an operation in H , reported with the number of distinct cases they represent.
Complete pipeline	A system that generates its own proposal. REVIEW: a draft plus one general-review call.
Controlled path	Extraction, proposal, and WITNESS admission; reply rendered from the execution log.
Strict correction completion	Allowed action and exactly the reference-specified whole state on the 11 correction cases.
Positive-change completion	Allowed action and exact whole state on the 28 positive cases (the prespecified primary endpoint).
Reference-violating state	Forbidden or extra write, or damage to a protected fact, as defined by the reference (all 48 cases).
Whole-state conformity	Executed state meeting every state requirement of the reference (all 48 cases).
Negative preservation	Preservation of the state on the 20 cases that require no change.

Re-evaluations of inference benchmarks, recommenders, and metric learning show how evaluation choices can change conclusions without a new model [11, 15, 28]. Counterfactual replay [37] and performance bounds under incomplete labels [27] provide related methodological context. Here, exposure complements power analysis [5] by bounding the endpoint difference possible on the saved outputs.

3 Task and Decision Boundary

A test case is $x = (h, T, M, \pi)$, consisting of the preceding context h ; an indexed sequence $T = (t_1, \dots, t_k)$ of alternative readings of one current utterance, which may coincide; an initial memory state M ; and a case-level retention permission π . The readings are simultaneous alternatives for the same speaker’s turn, not separate speakers or successive turns. A record has a persistent identifier and a fact tuple (subject, relation, value, time). The model proposes an action a , a reply u , and an ordered list of memory operations Q ; the application executes the admitted operations in one transaction and returns the state M' together with an execution log ρ .

Four operations are available: APPEND, CORRECT, HOLD, and no-write. A correction targets an existing record and replaces its value while keeping its identifier, whereas a hold leaves a proposed fact uncommitted. Retention permission is separate from factual support, because a clearly stated fact may still be disallowed for storage. Supplying the initial state lets the evaluation distinguish a new fact from an obsolete or protected one without having to infer it. For each case, the reference specifies the allowed actions, the required and forbidden tuples, the records that must remain unchanged, the obsolete versions to retire, and any records that must no longer appear; all of these requirements are checked on the executed state.

We assess required change and preservation together. On a positive case, at least one required fact is absent or an obsolete version must be retired, so holding every proposal fails. A correct target update can still fail whole-state conformity when it also appends a fact outside the requested edit. We score the action and persistent state against the reference and inspect reply–state discrepancies in selected traces; reply truthfulness is not a scored endpoint. In tables, “Viol.” abbreviates reference violations. Table 2 defines the terms and endpoints used throughout.

Table 3. Admission rules applied to the same proposal. Outside H , every rule keeps AGREE’s decision; retention permission and execution checks are shared.

Rule	Admission within H
AGREE	Keeps the conservative hold when shared extracted support is missing.
ALLOW	Admits every otherwise eligible correction in H .
LITERAL	Checks the existing target and the unchanged subject, relation, and time; requires the new value in every reading and current-turn evidence.
WITNESS	Additionally requires one affirmed, matching, target-linked witness per reading, supported by an exact source quotation that contains the value.

4 Correction Admission Rules

4.1 One Proposal and Four Admission Decisions

The controlled path makes two model calls. The first returns an extraction F : for each reading, the atomic facts that the reading supports, and, for each pair of a stored record and a reading, one correction witness with a polarity label (affirmed, negated, uncertain, or absent), a new value, and a verbatim quotation. The second call receives the test case together with F , marked as fallible, and generates the final proposal. Every admission rule receives the identical test case, extraction, and proposal, together with an isolated copy of M , so changing the rule requires no further model call (Figure 2).

AGREE applies to every proposed APPEND and CORRECT operation. It admits an operation when the exact subject–relation–value–time tuple is supported by the fact list of every reading or is already present in M ; otherwise, it changes the operation to HOLD. AGREE was the conservative gate of the original protocol, fixed before collection, and is meant to keep a fact out of memory when the readings disagree about it. Shared support requires that the extraction validly cover all readings. Because the test compares tuples exactly as extracted, and because the extraction schema allows a reading’s fact list to be empty, a correction can appear in the proposal and in the witnesses yet be missing from the intersection of the extracted fact lists. This is how the held set arises.

Let $H(x, F, Q)$ be the set of operation indices j for which AGREE holds a proposed CORRECT only because its support is not shared, although every reading is validly covered. All differences between AGREE and the other three rules lie within H ; outside H , every rule keeps AGREE’s decision (Table 3). ALLOW restores every correction in H and nothing else, so an APPEND that lacks shared support stays held under every rule. Two predicates define the evidence checks. For LITERAL ($L_j = 1$), the proposal must name one existing target; preserve its subject, relation, and time; change its value; and cite the current turn in both the fact and the operation. In addition, the replacement value must occur literally in every reading. For WITNESS ($W_j = 1$), the same target and source conditions must hold, and each reading must supply exactly one witness for the target record that is labeled affirmed, repeats the proposed subject, relation, time, and value, cites the current turn, and quotes a verbatim span of that reading containing the replacement value. The application verifies the structure and linkage of these model-generated witnesses, but the polarity label itself is taken from the model.

Two examples show how the checks differ. The sentence “the shop opened in 1993, not 1991” satisfies both checks for an opening-year record. By contrast, “My son works at Daehan Post” satisfies the literal check for an update of the speaker’s own workplace but should fail the witness check, which requires every reading to affirm the value for the targeted record: the affirmed statement concerns the son.

4.2 Containment and the Opportunity to Intervene

The admitted sets of the four rules are nested. Write A_X for the indices in H that rule X admits; then

$$A_{\text{WITNESS}} = \{j \in H : W_j = 1\} \subseteq A_{\text{LITERAL}} = \{j \in H : L_j = 1\} \subseteq A_{\text{ALLOW}} = H. \quad (1)$$

The first inclusion holds because a passing witness carries the required target and source linkage and an exact quotation that contains the value, so the value occurs in every reading; the second follows from the restriction to H . This ordering concerns admitted operations only. It need not carry over to final-state accuracy when a transaction is rejected, when operations interact, or when a witness is wrong.

AGREE and ALLOW are therefore the two ends of this ordering rather than competing designs: no rule that acts on H can admit more held corrections than ALLOW or fewer than AGREE. We compare LITERAL and WITNESS with both ends: with ALLOW for the reference violations they prevent, and with AGREE for the corrections they recover and the reference violations they add. Because AGREE reads model-extracted facts, its holds can originate in extraction rather than in disagreement between readings; the saved extraction shows which of the two occurred (Section 6.1). WITNESS and LITERAL decide differently only on $K = \{j \in H : L_j = 1, W_j = 0\}$, where a hold may either protect the correct state or keep an obsolete fact unnecessarily. When H is empty, all four rules must agree.

Admission and execution are deterministic given the saved case, extraction, and proposal. Replay therefore identifies effects conditional on these artifacts, rather than effects of regenerating a complete pipeline [17]. Exposure bounds the endpoint differences such a comparison can produce.

PROPOSITION 4.1 (EXPOSURE BOUND AND REFERENCE DEPENDENCE). *Let rules X and Y act on the same N outputs, and let E be the set of outputs that contain an operation on which X and Y can decide differently. For every endpoint indicator S computed on the executed state, (i) $|\bar{S}_X - \bar{S}_Y| \leq |E|/N$; (ii) the difference $\bar{S}_X - \bar{S}_Y$ depends on the reference only through the scores of the outputs in E ; and (iii) if a second reference changes S on at most C of the $2|E|$ output–rule pairs in E , the difference changes by at most C/N .*

PROOF. Outside E , the two rules admit the same operations, so the deterministic executor produces the same state and S takes the same value under X and Y . Hence $N(\bar{S}_X - \bar{S}_Y) = \sum_{o \in E} (S_X(o) - S_Y(o))$, which is at most $|E|$ in absolute value, involves the reference only through outputs in E , and changes with the reference by at most the number of its terms' scores that change. $\square$

For the four correction rules, E contains the outputs that reach H . For gating versus not gating, it also includes outputs with held appends or invalid extractions. With 11 exposed outputs among 144, an endpoint difference cannot exceed 7.6 percentage points; with zero exposure, the data do not test decisions on which the rules differ. Reference review must use this gate-log-based set, which can include outputs that tie under the current scores (Section 6.4).

4.3 Execution and Complete Pipelines

The application applies operations transactionally to a separate SQLite state for each rule and records read-after-write execution logs; the executor enforces retention permission and target existence for every rule. Because admission does not regenerate the action or the reply, differences in state outcomes between rules can be assigned to admission alone within this pipeline. For the same reason, a held write can coexist with a reply that asserts the held value (Online Appendix H), and the unchanged reply is not credited to admission. WITNESS uses the same two calls per output as AGREE, and ALLOW and LITERAL are computed from the saved proposals (Algorithm 1; Online Appendix C).




Fig. 2. Fixed-proposal architecture. Two generation calls produce one extraction and one proposal. Each admission rule is then applied to the same artifacts and executed on an isolated copy of the initial state.

Algorithm 1 Admission on a fixed proposal. Reply rendering is a separate downstream option.

Require: case $x = (h, T, M, \pi)$, extraction F , proposal (a, u, Q) , rule X (AGREE, ALLOW, LITERAL, or WITNESS)

Ensure: action a and reply u unchanged; executed state M' and log ρ

```

1: if  $Q$  fails schema validation then
2:   return  $(a, u, M, \emptyset)$  ▷ no write; scored as an invalid output
3: end if
4:  $Q_X \leftarrow \text{AGREEGATE}(x, F, Q)$  ▷ holds APPEND and CORRECT without shared support
5: for all  $j \in H(x, F, Q)$  do
6:   if  $X = \text{ALLOW}$  or  $(X = \text{LITERAL and } L_j = 1)$  or  $(X = \text{WITNESS and } W_j = 1)$  then
7:      $Q_{X,j} \leftarrow Q_j$  ▷ restore the proposed correction
8:   end if
9: end for
10:  $(M', \rho) \leftarrow \text{EXECUTE}(M, Q_X, \pi)$ 
11: return  $(a, u, M', \rho)$ 

```

The comparison pipeline, REVIEW, generates its own proposal from a draft and one general-review call, so it and the controlled path use the same number of calls. The original protocol also ran a single direct draft and one draft-feedback-refine cycle adapted from Self-Refine [24]; they share REVIEW's initial draft where the design permits and are reported in Online Appendices A and F. The prespecified primary test compares the controlled path, with its reply rendered from the execution log, with REVIEW. Rendering does not change any state endpoint, so the tables report the common state score under WITNESS (Online Appendix A).

5 Evaluation Design

5.1 Sources and Model Configurations

The main source contains 60 Korean test cases written by the authors with AI assistance: 12 development cases and 48 held-out cases spanning 12 problem families of four cases each (Online Appendix A, Table A3). Each held-out case was run three times per model configuration. A private reference, fixed before the main collection, specifies the requirements listed in Section 3, including the distinction between replacing a value and removing its record. The two

Table 4. Sources and their roles. Additional configurations reuse the main cases, and each source keeps its own denominators.

Source	Cases	Repetitions	Role
Main Korean cases	48	3	Fixed-proposal and complete-pipeline comparisons
DSTC2 test turns	48	1	Real recognition hypotheses; 41 callers
Contrast set	24	1	12 supported and 12 unsupported edits
Boundary probe	12	1	Six paired edits; no exposure
Public dialogue probe	16	1	10 eligible self-repairs; six unscored
Follow-up runs on main cases	5; 48	1	Revised extraction; learned verifier

co-first authors independently checked each of the 48 held-out references against its case. The collection script received the public cases but not the reference, and the configurations, prompts, admission rules, and evaluator were fixed in a frozen protocol before collection. We use Korean because widely used memory benchmarks are in English [26, 46] and because Korean writing yields several surface forms of one fact (word spacing varies in practice, and a person may be named by a kinship term, a given name, or both), which stresses exact tuple matching (Section 6.4).

All alternative readings of the main source are written text authored with AI assistance. Within a case, they can be identical, differ in surface form, or differ in value, person, or polarity; they play the role of competing recognition hypotheses but were not produced by a recognizer, so how often they reach the held set characterizes this source rather than any recognizer.

The original protocol planned a Holm-adjusted primary comparison for five model families; Gemini 2.5 Flash, Llama 3.1 8B, and Qwen3 14B (local 4-bit) were collected, and two families were not (Online Appendix C). Six later Gemini configurations reuse the 48 cases and were added after the earlier results had been inspected. The model panel comprises the 21 configurations that completed all 48 cases in all three repetitions with at least 95% valid final outputs in both REVIEW and the controlled path: seven Gemini, eight Claude, and three GPT configurations together with DeepSeek 3.2, GLM 5, and Qwen3 14B, for 3,024 case-repetition episodes. This validity criterion was set after collection and applies to complete pipelines, not semantic outcomes; of the 28 configurations with complete triplets, seven fall below it, among them Llama 3.1 8B (Online Appendix G). Model labels denote recorded provider or runtime configurations and do not imply matched parameter counts, precision, or reasoning budgets.

Table 4 lists the further sources. Three authored or adapted sources probe transfer and are reported in Online Appendices D–F: a 24-case contrast set written after the initial finding by a separate agent that saw neither the rule code nor earlier outputs, a 12-case boundary probe, and 16 selected MultiWOZ 2.4 turns [51]. A fourth source, 48 user turns from the DSTC2 test set, supplies real speech-recognition hypotheses with a reference derived from human-reviewed labels (Section 6.7). Three follow-up runs, collected after the results above had been inspected and fixed before their outputs, test the diagnoses (Online Appendix J): a revised extraction instruction on the held cases, a learned verifier (Gemini 3.1 Pro Preview) that approves or holds each APPEND and CORRECT in saved REVIEW proposals without seeing the extraction or the reference, and the DSTC2 replay.

5.2 Endpoints and the Variant-Tolerant Reference

The primary endpoint prespecified in the original study is positive-change completion: joint success of the action and the whole state on the 28 of the 48 main cases that require a new fact or the retirement of an obsolete version. A successful output must be valid, choose an allowed action, contain the required facts, omit forbidden and obsolete tuples, preserve protected facts, satisfy the required removals, and introduce no unlisted fact; provider failures and

invalid outputs stay in the denominator. Strict correction completion applies the same criterion to the 11 correction-eligible cases. The reference-violation endpoint flags forbidden or extra writes and damage to protected facts across all 48 cases. Successful positive changes and reference-violating states are always reported together, so that withholding more operations is not mistaken for improving the task.

The reference matches tuples exactly, so an unlisted fact counts as a violation even when it restates a required fact in another form. To measure how much each conclusion depends on this convention, we re-scored every saved output under a *variant-tolerant reference*. A tuple counts as a required or protected fact when its relation and time match and, after Unicode NFC, case folding, and whitespace removal, its subjects are equal, both denote the participant, or one contains the other, and its values are equal or one contains the other; a participant label is never matched to another person’s label. Forbidden and obsolete tuples stay exact, so the second reference can only credit outputs. It accepts 38 distinct pairs in 10 cases, all listed in Online Appendix K, and the scorer reproduces every archived score of the exact reference. It is a mechanical convention, not a semantic judgment, and no person judged the equivalences it accepts.

5.3 Paired Analysis

For configuration m and case i , the paired difference is $d_{mi} = \frac{1}{3} \sum_{r=1}^3 [Y_{mir}(\text{ctrl}) - Y_{mir}(\text{REVIEW})]$, where Y is the endpoint indicator for repetition r and ctrl denotes the controlled path; the effect estimate averages d_{mi} over cases. Uncertainty is reported as 95% percentile intervals from 5,000 case-cluster bootstrap draws that keep the repetitions, and for pooled estimates the configurations, of a case together (48 clusters for all-case endpoints, 28 for positive completion, and 11 for correction completion). The prespecified exact two-sided McNemar test counts a case as a success when it succeeds in at least two of three repetitions, and the original plan adjusted the five planned family tests with Holm’s method (Online Appendix A). Each original-protocol run made $48 \times 3 \times 6 = 864$ physical calls, of which each two-call pipeline used 288, without retries. Gemini 2.5 Flash used temperature 0.2, disabled thinking, and 2,048 output tokens; later configurations used their supported thinking settings, so the panel compares configurations as run rather than at equal compute (Online Appendix C). Counts pool configurations and repetitions, which are not independent replications of a case; distinct cases are reported alongside.

5.4 Confirmatory and Exploratory Analyses

Only the comparison of the controlled path with REVIEW was prespecified as a test; the study began as that test, and the replay analyses were added to explain its outcome (Table 5). Applying the gate to saved general-review proposals pairs each saved REVIEW proposal with the extraction from the same configuration, case, and repetition, applies one of the four rules, and executes the result with the frozen executor; we write REVIEW+WITNESS and so on. These conditions, two string-normalization controls, the variant audit, and the variant-tolerant re-scoring required no model call (Online Appendices I and K).

6 Results

Component replay varies admission while holding the extraction and proposal fixed; complete-pipeline comparisons also let the generated artifacts change. We report the first before the second.

6.1 Finding 1: Extraction Omissions Create Recoverable Holds

On Gemini 2.5 Flash, the agreement gate withheld 11 required corrections that the other three rules admitted. Changing admission from AGREE to LITERAL or WITNESS increases completed corrections from 22 to 33 of 33 observations,

Table 5. Status of the reported analyses. "Prespecified" does not imply public preregistration.

Analysis	Inferential role
Controlled path versus REVIEW (original families)	Prespecified primary test; five-slot Holm accounting retained.
AGREE versus WITNESS	Planned secondary contrast; descriptive.
ALLOW, LITERAL, hold traces, exposure; gate on REVIEW proposals	Post hoc, on saved artifacts.
Variant audit, variant-tolerant re-scoring, decomposition	Post hoc and mechanical; no model call.
Later Gemini configurations; 21-configuration panel	Exploratory reuse of the same 48 cases.
Further sources and follow-up runs	Diagnostic; fixed before their outputs.

Table 6. Comparisons on identical proposals in the seven Gemini configurations, in the order AGREE / ALLOW / LITERAL / WITNESS, under the exact reference. Strict correction uses 33 observations from 11 cases, and the reference-violation endpoint uses 144 observations from 48 cases. The H column reports exposed outputs and, in parentheses, distinct cases; by Proposition 4.1, no two rules can differ by more than this count. All counts combine three repetitions.

Configuration	Correction /33	Viol. /144	H outputs (cases)
Gemini 2.5 Flash	22 / 33 / 33 / 33	3 / 3 / 3 / 3	11 (5)
Gemini 2.5 Pro	31 / 31 / 31 / 31	3 / 9 / 4 / 3	6 (3)
Gemini 2.5 Flash-Lite	31 / 31 / 31 / 31	16 / 19 / 16 / 16	3 (1)
Gemini 3 Flash Preview	33 / 33 / 33 / 33	12 / 12 / 12 / 12	0 (0)
Gemini 3.1 Pro Preview	33 / 33 / 33 / 33	2 / 2 / 2 / 2	0 (0)
Gemini 3.5 Flash-Lite	32 / 32 / 32 / 32	7 / 12 / 7 / 7	5 (2)
Gemini 3.8 Flash	33 / 33 / 33 / 33	1 / 1 / 1 / 1	0 (0)

while reference-violating states remain at 3 of 144 under every rule (Table 6). The 11 recovered corrections come from five of the 11 eligible cases, and eight are repetitions of self-repair cases. The paired case-mean gain is 33.3 percentage points (95% case-cluster interval 12.1 to 57.6), and the case-level McNemar test, with four discordant cases all favoring WITNESS, gives $p = .125$; four discordant cases are too few to resolve the gain, so we report it as a count of cases (Section 8). Because ALLOW recovers the same corrections, releasing the hold is sufficient for recovery on these proposals.

In all 11 held observations, both extracted fact lists were empty, although the correction appeared as two affirmed witnesses matching the proposal. In the other 22 observations, both fact lists contained the corrected tuple. Thus the holds arose from extraction omissions, not conflicting readings. The contrast set contains five further supported corrections with the same omission pattern, all recovered by the three releasing rules. Only one of the 17 traced holds reflects disagreement: one reading negates the new residence (Online Appendix D).

A follow-up re-extracted one held repetition from each of the five cases under the original instruction and a revision that also lists affirmed corrections among the extracted facts (Online Appendix J). The saved proposal and AGREE gate were unchanged. The original instruction completed one of five corrections; the revision completed all five, with no reference violation in either arm. CT01 recovered spontaneously under the original instruction and is not an effect of the revision. On Gemini 3.8 Flash, which had zero exposure, both instructions produced identical outcomes on all 48 cases. The revised extractor was therefore sufficient to remove the selected Flash holds; this single follow-up does not establish its effect on new cases.

6.2 Finding 2: Zero Exposure Explains Many Ties

Gemini 3 Flash Preview, 3.1 Pro Preview, and 3.8 Flash produce no output that reaches H (Table 6). Their ties are instances of the opportunity confound, as is the 12-case boundary probe, in which no output reached H and every rule produced the same states (Online Appendix E). Across the 21 panel configurations, the saved gate decisions on controlled-path proposals show no exposed output on 12 configurations and 49 exposed outputs, from 12 cases, on the other nine (Online Appendix G). Because every configuration answers the same 48 cases, the range of exposure, from 0 to 11 of 144 outputs, shows that whether this benchmark can separate admission rules depends on the generator; on general-review proposals, 16 of the 21 configurations have no exposed output (Online Appendix I). Exposure therefore has to be reported per configuration, not per benchmark. By Proposition 4.1, no panel-wide comparison among the rules on these proposals can move an endpoint by more than 49 of 3,024 outputs, or 1.6 percentage points. The same count states this benchmark’s main limitation: the rules can differ on only 12 of its 48 cases (Section 8). Both evidence checks release 12 of the 49 holds, each completing a required correction (11 on Gemini 2.5 Flash and one on DeepSeek 3.2); LITERAL alone releases two more, both reference violations; and all three rules keep the other 35, all in cases that require no change, where the held state conforms to the reference. On most exposed outputs, the agreement gate thus withheld writes that the reference rejects, as designed.

Within the exposed outputs, a literal check supplies most of the protection that the evidence checks add. Relative to ALLOW, WITNESS avoids 14 reference-violating state observations and LITERAL 13, all on Gemini 2.5 Pro, 2.5 Flash-Lite, and 3.5 Flash-Lite and from three distinct cases, while both keep ALLOW’s strict-correction count on every configuration (Table 6). On the saved Qwen3 14B proposals, ALLOW produces 29 reference-violating states out of 144 and LITERAL and WITNESS 20 each, with strict correction at 33 of 33. Compared with the conservative end, WITNESS matches AGREE’s reference-violation count on every configuration while recovering the 11 Flash corrections.

6.3 Target-Linked Evidence Blocked Person and Polarity Errors

The one reference-violating state that the literal check admits on Gemini 2.5 Pro illustrates what the witness check adds (Figure 3). The utterance describes a son’s workplace, and the same proposal both appends the son’s workplace and corrects the participant’s stored workplace. Because the replacement value occurs in both readings, LITERAL, like ALLOW, admits the second update; the witness for the participant’s record is not affirmed, so WITNESS holds it and keeps the participant’s workplace. Target-linked witnesses can thus block a value that appears in every reading but belongs to another person. The contrast set supplies a second situation, a value that one reading negates (Section 6.1), and the third is an uncertain report about a travel destination with a request not to change it (CT11), which LITERAL admits and WITNESS holds on GLM 5 controlled-path proposals and on saved Qwen3 14B review proposals. Literal occurrence cannot tell whose fact a value is or whether a reading affirms it; binding the evidence to the target record can, and in these four observations it changed the executed state. Section 6.7 adds a fifth, a negated value among real recognition hypotheses.

6.4 Finding 3: Reference Conventions Change the Gating Trade-off

Applying the gate to saved REVIEW proposals more than halves reference-violating states under the exact reference, from 146 to 64 with WITNESS, at a cost of 40 positive changes (Table 7). The two bounds separate the parts of the gate: REVIEW+AGREE already reaches 64, and REVIEW+ALLOW 75. Counting decisions by operation locates the difference (Online Appendix I). All four rules treat APPEND identically and hold 108 of 1,087 proposed appends; these holds




Fig. 3. Target binding changes the executed state. LITERAL accepts a workplace value that appears in both readings, whereas WITNESS rejects assigning the son's workplace to the participant. A correct-looking reply does not by itself guarantee a correct memory operation.

Table 7. The gate on saved REVIEW proposals under the exact and the variant-tolerant reference (21 configurations, 48 cases, three repetitions). Strict correction counts (690, 676, 688, 688, and 688 of 693) are the same under both references. The last column counts held decisions among 1,087 proposed APPEND and 762 proposed CORRECT operations.

Saved path	Viol. /3,024 ↓		Positive /1,764 ↑		Held APPEND / CORRECT
	Exact	Variant-tol.	Exact	Variant-tol.	
REVIEW	146	44	1,623	1,725	–
REVIEW + AGREE	64	14	1,571	1,621	108 / 23
REVIEW + ALLOW	75	25	1,583	1,633	108 / 0
REVIEW + LITERAL	65	15	1,583	1,633	108 / 10
REVIEW + WITNESS	64	14	1,583	1,633	108 / 11

account for 71 of the 82 avoided reference violations and 36 of the 40 lost positive changes, and the other four losses come from invalid extractions. AGREE holds 23 of 762 proposed corrections, on five configurations; ALLOW releases all of them, LITERAL 13, and WITNESS 12, and each release shared by these rules completes a required correction. On these proposals, the gate acts mainly as an append policy.

The avoided violations are not all semantic errors. Comparing each held tuple with the states that conform to the reference for the same case, 13 of the 82 involve a forbidden tuple, and 52 of the remaining 69 restate a required fact with different spacing (6), subject label (29; e.g., “어머니 순자” (“mother Sunja”) for “순자” (“Sunja”), or value form (17; e.g., “자전거 타기” (“riding a bicycle”) for “자전거” (“bicycle”). In these 52 observations the required fact is absent with or without the gate, so the hold lowers the violation count without changing completion.

The variant-tolerant reference measures how much this matters (Table 7). Under it, REVIEW produces 44 reference-violating states instead of 146 and 1,725 positive changes instead of 1,623: in 102 of its violating states, the only violation was a required fact written in another form. The gate then avoids 30 violations instead of 82 (paired difference 0.99 rather than 2.71 percentage points over 3,024 observations; case-cluster interval 0.40 to 1.69), and its cost rises from 40 to 92 positive changes (−5.22 points over 1,764, interval −8.22 to −2.49), because the appends it holds are now credited as required facts; whole-state conformity falls by 79 observations rather than 27. The reference therefore changes the apparent balance between protection and missed changes. Neither convention has independent human validation: the exact reference lists one form per fact, while the variant-tolerant reference accepts mechanically defined equivalences (Section 8).

Every output-level difference among the replayed correction rules is unchanged under the two references. Nonzero differences occur on 36 controlled-path and 23 review-proposal outputs; these score-based counts must not be substituted for the gate-log exposure counts of 49 and 23. The gap includes missing ALLOW replays and an exposed output whose scores tie (Online Appendices G and K). For REVIEW versus REVIEW+WITNESS, the reference changes the difference on 52 of the 127 outputs whose endpoints differ. Normalizing strings inside the gate changes no executed state on the main panel, whereas case differences affect the public task-dialogue probe (7 completed self-repairs under ALLOW versus 4 under the evidence checks, out of 10 eligible; Online Appendix F).

6.5 The Complete-Pipeline Difference Is Two Opposite Terms

Under the exact reference, complete-pipeline differences split evenly across configurations. Among the 21 panel configurations, the controlled path increases positive-change completion relative to general review on 9, decreases it on 9, and ties on 3, and reference violations decrease on 12, increase on 7, and tie on 2; 11 configurations change direction across repetitions (Online Appendix G). The prespecified primary comparison on the original families found no completion advantage. On Gemini 2.5 Flash, REVIEW completes 82 of 84 positive-change observations and the controlled path 80 (difference −2.4 percentage points, 95% interval −6.0 to 0.0; McNemar $p = 1.00$); on Qwen3 14B the difference is −8.3 points (−23.8 to 6.0); and Llama 3.1 8B, outside the panel because its controlled path produced 68 invalid final outputs, completes 6 of 84 against 29 (−27.4 points, −44.0 to −10.7; Holm-adjusted $p = .0586$).

An intermediate replay condition decomposes the observed difference: REVIEW+WITNESS applies the controlled path’s extraction and gate to the saved review proposal. The complete-pipeline difference equals the controlled path minus this intermediate condition (proposal term), plus the intermediate condition minus REVIEW (gate term; Table 8). This is a conditional decomposition through the chosen intermediate condition; it does not estimate a full factorial interaction. Under the exact reference, the controlled path’s proposals add 67 reference violations and 17 positive changes relative to review proposals under the same gate (violations +2.22 points, interval 1.03 to 3.67), and the gate removes 82 violations and 40 positive changes; the two terms nearly cancel, leaving −15 and −23. General review also completes all 33 strict corrections on Gemini 2.5 Flash, so the fixed-proposal gain has nothing to recover there. Under the variant-tolerant reference, the proposal term is +40 violations and +44 positive changes and the gate term −30 and −92; the controlled path then completes fewer positive changes on 13 configurations and more on one (pooled −2.72 points, −4.59 to −1.02), and violations decrease on 6 configurations and increase on 6. The direction of the complete-pipeline comparison thus depends on the reference, and the decomposition shows why: the gate’s append holds, not the extraction-informed proposals, carry the completion cost.

Table 8. Exact decomposition of the complete-pipeline difference, controlled path minus REVIEW, pooled over 21 configurations, 48 cases, and three repetitions (counts of observations; reference violations over 3,024, positive changes over 1,764). The proposal term compares two proposals under the same extraction and WITNESS gate; the gate terms attribute each paired change to the held operation, and no observation changed both.

Term	Reference violations		Positive changes	
	Exact	Variant-tol.	Exact	Variant-tol.
Proposal: controlled path – REVIEW+WITNESS	+67	+40	+17	+44
Gate: held APPEND	–71	–19	–36	–88
Gate: held CORRECT	–11	–11	0	0
Gate: invalid extraction	0	0	–4	–4
Complete pipeline: controlled path – REVIEW	–15	+10	–23	–48

Table 9. Admission rules on real recognition hypotheses (DSTC2, Gemini 3.1 Pro Preview, one run). Goal conformity requires the exact food, area, and price-range goal; required changes count the 36 turns whose goal changes; a new wrong write introduces a value absent from the official goal.

Rule	Goal conformity /48 ↑	Required changes /36 ↑	New wrong writes /48 ↓
Ungated proposal	33	23	3
ALLOW	33	23	3
AGREE	18	6	0
LITERAL	17	6	1
WITNESS	18	6	0

6.6 Replay Locates Where a Learned Verifier Acts

A learned verifier [57] replaced the gate on one saved REVIEW proposal per case for three configurations (144 proposals; Online Appendix J). On Gemini 2.5 Flash it matches REVIEW+WITNESS on every endpoint, including five more completed corrections than REVIEW+AGREE; these are the extraction-induced holds of Finding 1. Its other differences occur outside H : it admits a required append on Gemini 3.8 Flash and blocks two writes on Qwen3 14B, reducing reference violations from three to one. All paired intervals against REVIEW+WITNESS include zero. These observations locate the verifier’s decisions without establishing a ranking. The 144 calls averaged 4.39 seconds and about 850 tokens each.

6.7 On Real Recognition Hypotheses, Replay Finds a Different Cause

Unlike the written readings of the other sources, the DSTC2 source supplies hypotheses that a speech recognizer produced [16]. One user turn was selected from each of 48 test dialogues (41 callers) by a hash-ranked rule fixed before any output: 24 goal revisions, 12 denials without an explicit replacement, and 12 turns with recognition errors, a stratum selected with the labels and reported separately. The readings are the 433 distinct non-empty hypotheses of the live 10-best lists, in rank order and without likelihood weights. The initial memory is the previous official goal for food, area, and price range; the reference is the current official goal, derived from human-reviewed semantic labels and matched exactly; and the current transcript and labels were withheld from the model. The two-call pipeline, with the three slots in place of personal relations and the gate unchanged, ran once on Gemini 3.1 Pro Preview, all 48 outputs valid (Table 9); a partial Gemini 3.8 Flash run is reported in Online Appendix J.

Table 10. Implications of the three findings. Only the extraction revision was tested; the other rows identify follow-up evaluations.

Finding	Diagnostic check	Implication
Extraction omitted supported corrections	Compare with ALLOW; inspect saved fact lists	Revise extraction; test the fixed revision on new cases
Rules tied with zero exposure	Count outputs reaching H before scoring	Add cases that distinguish the rules; report the enriched sampling design
Reference conventions changed the gating trade-off	Re-score the same executed states	Validate fact equivalences before judging the append policy

Exposure is high on this source: AGREE held a proposed correction in 20 of the 48 outputs, against at most 11 of 144 for any main-source configuration. The trace also returns a different attribution. None of the 20 holds shows the pattern of Section 6.1, affirmed witnesses in every reading with the tuple missing from the fact lists; in ten, different hypotheses yield different extracted values, and in the other ten the value is supported by only some hypotheses or omitted from some. Releasing every hold (ALLOW) completes 17 required changes and introduces 3 wrong values. Neither evidence check releases any of the 17, because both require the value in every hypothesis, and LITERAL releases one hold, a wrong write. Relative to ALLOW, WITNESS lowers goal conformity by 31.25 percentage points (95% caller-cluster interval 16.3 to 45.8) and required-change completion by 47.2 points (32.3 to 62.9), and it lowers turns with a new wrong write by 6.25 points (0.0 to 13.5).

Two turns show the trade-off. In DSTC2_45, the caller says “i dont want chinese i want thai food”; the stored goal is already thai, but every hypothesis contains “chinese,” and the proposal changes the goal to chinese. LITERAL admits the change because the value occurs in every hypothesis, whereas WITNESS, like AGREE, holds it for lack of affirmed, target-linked witnesses: the polarity error of Section 6.3 on a real recognizer. In DSTC2_07, “how about vietnamese food” appears in some hypotheses as “how about you need food”; ALLOW admits the correct change, and the three strict rules hold it. On this source, agreement does what it was designed for, keeping values on which the hypotheses disagree out of memory, and replay prices it: requiring support in every reading suits written alternatives better than n-best lists.

7 Discussion

7.1 Interpreting Rule Comparisons

The three findings call for different checks. An ALLOW comparison and the saved extraction can distinguish evidence-based selection from recovery of upstream omissions. Exposure determines whether a tie tests competing decisions. Re-scoring tests whether a reported benefit depends on the reference’s representation of facts. These checks address different questions, and agreement under two mechanical references is not a substitute for human validation.

7.2 From Diagnosis to Action

Table 10 links the three findings to their implications. Only the extraction change was tested, once on five selected cases. The other actions remain proposals for further evaluation.

The DSTC2 result sets a boundary on the first diagnosis: there, holds arose from disagreement between recognizer hypotheses, not the extraction pattern found in the main source. Majority or likelihood-weighted support is a candidate intervention for that setting, with thresholds fixed on development data.

7.3 Replay Cost and Pipeline Resources

Replaying a deterministic rule requires no new model call, but deploying the hybrid gate adds an extraction stage to the two-call REVIEW pipeline. On Gemini 2.5 Flash, mean summed stage time rises from 3.00 to 4.92 seconds (64%); this is not a serving-latency measurement. The benefit to weigh against that cost depends on the reference (Finding 3). Larger memories would also require scoping witnesses to the records targeted by a proposal.

8 Limitations and Future Work

The study uses 48 authored Korean cases and 48 DSTC2 turns. In the main source, only 12 distinct cases reach H , five account for recovered corrections, and three for the literal check’s protection; the witness-only holds outside DSTC2 comprise four observations from three situations. Model configurations and repetitions increase observations, not independent cases. All replays use our own extraction-proposal pipeline, and the extraction intervention was tested once on selected cases. The results establish mechanisms in these data, not their prevalence in other writers or a general ranking of admission rules. The only confirmatory comparison found no completion advantage; under the exploratory variant-tolerant reference, the controlled path completes fewer positive changes than general review.

The authors created the main cases and references with AI assistance, and both co-first authors independently checked the held-out references. No external annotator checked them, no agreement statistic was computed, and no person validated the mechanical variant equivalences. Only DSTC2 supplies an externally labeled reference. Witness polarity remains a model judgment despite structural and source-link checks. We do not assess semantic safety, user utility, later retrieval, or longitudinal interaction.

The next priorities are independent human validation and replication on externally authored cases and another memory writer. For the four correction rules on the main panel, review should start from the 49 controlled-path and 23 review-proposal exposed outputs, after completing the missing ALLOW replays on 14 outputs, together with the 38 accepted variant pairs. Score ties alone do not justify excluding an output; gating-versus-not-gating and learned-verifier comparisons require their own review sets. Exposure can guide additional diagnostic sampling, which must be distinguished from prevalence estimation. Extraction revisions and admission thresholds should be fixed on development data before testing on new cases.

9 Conclusion

Replay of saved memory-write proposals identified three mechanisms: extraction omissions created recoverable holds, zero exposure explained many rule ties, and reference conventions changed the apparent benefit of gating. The practical consequence is to inspect the source of held writes, report where rules can differ, and test reference sensitivity before attributing a gain to admission. These findings are supported within the tested pipeline; independent human judgments and evaluation on new cases and writers are needed to establish their broader use.

Data, Code, and AI Use

The archive at <https://anonymous.4open.science/r/memory-admission-9E37/> (artifact version paper-artifact-v8) contains the Korean test cases, references, prompts, raw outputs and failures, executed states, the evaluator, and the analysis code, including the variant-tolerant re-scoring, with provenance hashes; for DSTC2, which is used under its own terms and not redistributed, it contains the selected turn identifiers and the code that reconstructs the inputs. The new experiments involved no human participants. ChatGPT and Codex assisted with language editing, manuscript

organization, and analysis code; their use in drafting the test cases is described in Section 5.1. The authors reviewed all AI-assisted content and take full responsibility for it. English examples are translations of the Korean inputs that define the experiments.

References

- [1] Vaibhav Adlakha, Parishad BehnamGhader, Xing Han Lu, Nicholas Meade, and Siva Reddy. 2024. Evaluating Correctness and Faithfulness of Instruction-Following Models for Question Answering. *Transactions of the Association for Computational Linguistics* 12 (2024), 681–699. doi:10.1162/tac1_a_00667
- [2] Anirudh Ajith, Chris Pan, Mengzhou Xia, Ameet Deshpande, and Karthik Narasimhan. 2024. InstructEval: Systematic Evaluation of Instruction Selection Methods. In *Findings of the Association for Computational Linguistics: NAACL 2024*. Association for Computational Linguistics, 4336–4350. doi:10.18653/v1/2024.findings-naacl.270
- [3] Jacob Andreas, John Bufe, David Burkett, Charles Chen, Josh Clausman, Jean Crawford, Kate Crim, Jordan DeLoach, Leah Dörner, Jason Eisner, et al. 2020. Task-Oriented Dialogue as Dataflow Synthesis. *Transactions of the Association for Computational Linguistics* 8 (2020), 556–571. doi:10.1162/tac1_a_00333
- [4] Sanghwan Bae, Donghyun Kwak, Soyoung Kang, Min Young Lee, Sungdong Kim, Yuin Jeong, Hyeri Kim, Sang-Woo Lee, Woomyoung Park, and Nako Sung. 2022. Keep Me Updated! Memory Management in Long-term Conversations. In *Findings of the Association for Computational Linguistics: EMNLP 2022*. Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 3769–3787. doi:10.18653/v1/2022.findings-emnlp.276
- [5] Dallas Card, Peter Henderson, Urvashi Khandelwal, Robin Jia, Kyle Mahowald, and Dan Jurafsky. 2020. With Little Power Comes Great Responsibility. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 9263–9274. doi:10.18653/v1/2020.emnlp-main.745
- [6] Hanxiang Chao, Yihan Bai, Rui Sheng, Tianle Li, and Yushi Sun. 2026. STALE: Can LLM Agents Know When Their Memories Are No Longer Valid? arXiv:2605.06527 doi:10.48550/arXiv.2605.06527
- [7] Zhaorun Chen, Zhen Xiang, Chaowei Xiao, Dawn Song, and Bo Li. 2024. AgentPoison: Red-teaming LLM Agents via Poisoning Memory or Knowledge Bases. In *Advances in Neural Information Processing Systems*, Vol. 37. 130185–130213. doi:10.52202/079017-4136
- [8] Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. 2025. Mem0: Building Production-Ready AI Agents with Scalable Long-Term Memory. arXiv:2504.19413 doi:10.48550/arXiv.2504.19413
- [9] Shen Dong, Shaochen Xu, Pengfei He, Yige Li, Jiliang Tang, Tianming Liu, Hui Liu, and Zhen Xiang. 2025. Memory Injection Attacks on LLM Agents via Query-Only Interaction. In *Advances in Neural Information Processing Systems*, Vol. 38. <https://openreview.net/forum?id=i7J62t2wtV>
- [10] Yijiang River Dong, Tiancheng Hu, Zheng Hui, Caiqi Zhang, Ivan Vulić, Andreea Bobu, and Nigel Collier. 2026. Value of Information: A Framework for Human-Agent Communication. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 42879–42896. doi:10.18653/v1/2026.acl-long.1987
- [11] Maurizio Ferrari Dacrema, Paolo Cremonesi, and Dietmar Jannach. 2019. Are We Really Making Much Progress? A Worrying Analysis of Recent Neural Recommendation Approaches. In *Proceedings of the 13th ACM Conference on Recommender Systems*. ACM, 101–109. doi:10.1145/3298689.3347058
- [12] Felix Gervits, Gordon Briggs, Antonio Roque, Genki A. Kadomatsu, Dean Thurston, Matthias Scheutz, and Matthew Marge. 2021. Decision-Theoretic Question Generation for Situated Reference Resolution: An Empirical Study and Computational Model. In *Proceedings of the 2021 International Conference on Multimodal Interaction*. ACM, 150–158. doi:10.1145/3462244.3479925
- [13] Shengyue Guan, Jindong Wang, Jiang Bian, Bin Zhu, Jian-Guang Lou, and Haoyi Xiong. 2026. Evaluating LLM-based Agents for Multi-turn Conversations: A Survey. *ACM Transactions on Intelligent Systems and Technology* 17, 4, Article 85 (2026), 40 pages. doi:10.1145/3793671
- [14] Jinyu Guo, Zhaokun Wang, Jingwen Pu, Wenhong Tian, Guiduo Duan, and Guangchun Luo. 2025. Multi-Perspective Dialogue Non-Quota Selection with loss monitoring for dialogue state tracking. *Expert Systems with Applications* 283 (2025), 127516. doi:10.1016/j.eswa.2025.127516
- [15] Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel Bowman, and Noah A. Smith. 2018. Annotation Artifacts in Natural Language Inference Data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*. Association for Computational Linguistics, 107–112. doi:10.18653/v1/N18-2017
- [16] Matthew Henderson, Blaise Thomson, and Jason D. Williams. 2014. The Second Dialog State Tracking Challenge. In *Proceedings of the 15th Annual Meeting of the Special Interest Group on Discourse and Dialogue (SIGDIAL)*. Association for Computational Linguistics, Philadelphia, PA, USA, 263–272. doi:10.3115/v1/W14-4337
- [17] Paul W. Holland. 1986. Statistics and Causal Inference. *J. Amer. Statist. Assoc.* 81, 396 (1986), 945–960. doi:10.1080/01621459.1986.10478354
- [18] Or Honovich, Roei Aharoni, Jonathan Herzig, Hagai Taitelbaum, Doron Kukliansy, Vered Cohen, Thomas Scialom, Idan Szpektor, Avinatan Hassidim, and Yossi Matias. 2022. TRUE: Re-evaluating Factual Consistency Evaluation. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, Seattle, United States, 3905–3920. doi:10.18653/v1/2022.naacl-main.287
- [19] Sungdong Kim, Sohee Yang, Gyuwan Kim, and Sang-Woo Lee. 2020. Efficient Dialogue State Tracking by Selectively Overwriting Memory. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, 567–582. doi:10.1162/tac1_a_00333

Manuscript submitted to ACM

- .18653/v1/2020.acl-main.53
- [20] Hao Li, Chenghao Yang, An Zhang, Yang Deng, Xiang Wang, and Tat-Seng Chua. 2025. Hello Again! LLM-powered Personalized Agent for Long-term Dialogue. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics, 5259–5276. doi:10.18653/v1/2025.naacl-long.272
- [21] Hongjun Liu and Chen Zhao. 2026. Memory Is a Derivation: The Distributed-Evidence Paradox in Long-Term Agents. arXiv:2609.36130 doi:10.48550/arXiv.2609.36130
- [22] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranajpe, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics* 12 (2024), 157–173. doi:10.1162/tac_l_a_00638
- [23] Zhanghui Liu, Chaoxiong Zhou, Yuzhong Chen, and Zeping Zheng. 2026. Few-shot dialogue state tracking via context ranking and dynamic contrastive decoding. *Expert Systems with Applications* 300 (2026), 130197. doi:10.1016/j.eswa.2025.130197
- [24] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. 2023. Self-Refine: Iterative Refinement with Self-Feedback. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., 46534–46594. doi:10.52202/075280-2019
- [25] Chris Madge, Matthew Purver, and Massimo Poesio. 2025. Referential ambiguity and clarification requests: comparing human and LLM behaviour. In *Proceedings of the Eighth Workshop on Computational Models of Reference, Anaphora and Coreference*. Association for Computational Linguistics, 1–11. doi:10.18653/v1/2025.crac-1.1
- [26] Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. 2024. Evaluating Very Long-Term Conversational Memory of LLM Agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 13851–13870. doi:10.18653/v1/2024.acl-long.747
- [27] Felipe Maia Polo, Subha Maity, Mikhail Yurochkin, Moulinath Banerjee, and Yuekai Sun. 2024. Weak Supervision Performance Evaluation via Partial Identification. In *Advances in Neural Information Processing Systems*, Vol. 37.
- [28] Kevin Musgrave, Serge Belongie, and Ser-Nam Lim. 2020. A Metric Learning Reality Check. In *Computer Vision – ECCV 2020 (Lecture Notes in Computer Science, Vol. 12370)*. Springer, 681–699. doi:10.1007/978-3-030-58595-2_41
- [29] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. 2023. MemGPT: Towards LLMs as Operating Systems. arXiv:2310.08560 doi:10.48550/arXiv.2310.08560
- [30] Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, Xufang Luo, Hao Cheng, Dongsheng Li, Yuqing Yang, Chin-Yew Lin, H. Vicky Zhao, Lili Qiu, and Jianfeng Gao. 2025. SeCom: On Memory Construction and Retrieval for Personalized Conversational Agents. In *International Conference on Learning Representations*, Vol. 2025. 91851–91885. https://proceedings.iclr.cc/paper_files/paper/2025/hash/e56f394bbd4f0ec81393d767caa5a31b-Abstract-Conference.html
- [31] Joon Sung Park, Joseph O’Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative Agents: Interactive Simulacra of Human Behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. ACM, 1–22. doi:10.1145/3586183.3606763
- [32] Sudha Rao and Hal Daumé III. 2018. Learning to Ask Good Questions: Ranking Clarification Questions using Neural Expected Value of Perfect Information. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 2737–2746. doi:10.18653/v1/P18-1255
- [33] Hannah Rashkin, Vitaly Nikolaev, Matthew Lamm, Lora Aroyo, Michael Collins, Dipanjan Das, Slav Petrov, Gaurav Singh Tomar, Iulia Turc, and David Reitter. 2023. Measuring Attribution in Natural Language Generation Models. *Computational Linguistics* 49, 4 (2023), 777–840. doi:10.1162/coli_a_00486
- [34] Preston Rasmussen, Pavlo Paliychuk, Travis Beauvais, Jack Ryan, and Daniel Chalef. 2025. Zep: A Temporal Knowledge Graph Architecture for Agent Memory. arXiv:2501.13956 doi:10.48550/arXiv.2501.13956
- [35] Wout Schellaert, Fernando Martínez-Plumed, and José Hernández-Orallo. 2025. Analysing the Predictability of Language Model Performance. *ACM Transactions on Intelligent Systems and Technology* 16, 2, Article 30 (2025), 26 pages. doi:10.1145/3706118
- [36] Ivan Sekulić, Mohammad Aliannejadi, and Fabio Crestani. 2024. Analysing Utterances in LLM-Based User Simulation for Conversational Search. *ACM Transactions on Intelligent Systems and Technology* 15, 3, Article 62 (2024), 22 pages. doi:10.1145/3650041
- [37] Jaineet Shah. 2026. Causal Agent Replay: Counterfactual Attribution for LLM-Agent Failures. arXiv:2606.08275 doi:10.48550/arXiv.2606.08275
- [38] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language Agents with Verbal Reinforcement Learning. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., 8634–8652. doi:10.52202/075280-0377
- [39] Leila Tavakoli, Johanne R. Trippas, Hamed Zamani, Falk Scholer, and Mark Sanderson. 2025. Online and Offline Evaluation in Search Clarification. *ACM Transactions on Information Systems* 43, 1, Article 2 (2025), 30 pages. doi:10.1145/3681786
- [40] Dechuan Teng, Zheng Yan, and Wanxiang Che. 2026. Program of actions (PoA): Interactive semantic action parser as more powerful dialogue state tracker through code generation. *Expert Systems with Applications* 301 (2026), 130481. doi:10.1016/j.eswa.2025.130481
- [41] Alberto Testoni and Raquel Fernández. 2024. Asking the Right Question at the Right Time: Human and Model Uncertainty Guidance to Ask Clarification Questions. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 258–275. doi:10.18653/v1/2024.eacl-long.16

- [42] Xin Tian, Liankai Huang, Yingzhan Lin, Siqi Bao, Huang He, Yunyi Yang, Hua Wu, Fan Wang, and Shuqi Sun. 2021. Amendable Generation for Dialogue State Tracking. In *Proceedings of the 3rd Workshop on Natural Language Processing for Conversational AI*. Association for Computational Linguistics, 80–92. doi:10.18653/v1/2021.nlp4convai-1.8
- [43] Bo-Hsiang Tseng, Shruti Bhargava, Jiarui Lu, Joel Ruben Antony Moniz, Dhivya Piraviperumal, Lin Li, and Hong Yu. 2021. CREAD: Combined Resolution of Ellipses and Anaphora in Dialogues. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, 3390–3406. doi:10.18653/v1/2021.naacl-main.265
- [44] Qingyue Wang, Yanhe Fu, Yanan Cao, Shuai Wang, Zhiliang Tian, and Liang Ding. 2025. Recursively summarizing enables long-term dialogue memory in large language models. *Neurocomputing* 639 (2025), 130193. doi:10.1016/j.neucom.2025.130193
- [45] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *The Eleventh International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=1PL1NIMMrw>
- [46] Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. 2025. LongMemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory. In *International Conference on Learning Representations*, Vol. 2025. 86809–86836. https://proceedings.iclr.cc/paper_files/paper/2025/hash/d813d324dbf0598bbdc9c8e79740ed01-Abstract-Conference.html
- [47] Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024. Knowledge Conflicts for LLMs: A Survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Miami, Florida, USA, 8541–8565. doi:10.18653/v1/2024.emnlp-main.486
- [48] Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. 2025. A-Mem: Agentic Memory for LLM Agents. In *Advances in Neural Information Processing Systems*, Vol. 38. Curran Associates, Inc., Red Hook, NY, USA, 17577–17604. doi:10.52202/085713-0593
- [49] Sikuan Yan, Xiufeng Yang, Zuchao Huang, Ercong Nie, Zifeng Ding, Zonggen Li, Xiaowen Ma, Jinhe Bi, Kristian Kersting, Jeff Z. Pan, Hinrich Schütze, Volker Tresp, and Yunpu Ma. 2026. Memory-R1: Enhancing Large Language Model Agents to Manage and Utilize Memories via Reinforcement Learning. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, San Diego, California, United States, 12805–12825. doi:10.18653/v1/2026.acl-long.583
- [50] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2025. τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. In *International Conference on Learning Representations (ICLR)*. 9965–10017. https://proceedings.iclr.cc/paper_files/paper/2025/hash/1b126cc38b8638e07bef37e7b2bb72bf-Abstract-Conference.html
- [51] Fanghua Ye, Jarana Manotumruksa, and Emine Yilmaz. 2022. MultiWOZ 2.4: A Multi-Domain Task-Oriented Dialogue Dataset with Essential Annotation Corrections to Improve State Tracking Evaluation. In *Proceedings of the 23rd Annual Meeting of the Special Interest Group on Discourse and Dialogue*. Association for Computational Linguistics, Edinburgh, UK, 351–360. doi:10.18653/v1/2022.sigdial-1.34
- [52] Steve Young, Milica Gašić, Blaise Thomson, and Jason D. Williams. 2013. POMDP-Based Statistical Spoken Dialog Systems: A Review. *Proc. IEEE* 101, 5 (2013), 1160–1179. doi:10.1109/JPROC.2012.2225812
- [53] Xiaoxue Zang, Abhinav Rastogi, Srinivas Sunkara, Raghav Gupta, Jianguo Zhang, and Jindong Chen. 2020. MultiWOZ 2.2: A Dialogue Dataset with Additional Annotation Corrections and State Tracking Baselines. In *Proceedings of the 2nd Workshop on Natural Language Processing for Conversational AI*. Association for Computational Linguistics, 109–117. doi:10.18653/v1/2020.nlp4convai-1.13
- [54] Guilin Zhang, Wei Jiang, Xiejiashan Wang, Aisha Behr, Kai Zhao, Jeffrey Friedman, Xu Chu, and Amine Anoun. 2026. Adaptive Memory Admission Control for LLM Agents. arXiv:2603.04549 doi:10.48550/arXiv.2603.04549
- [55] Michael JQ Zhang and Eunsol Choi. 2025. Clarify When Necessary: Resolving Ambiguity Through Interaction with LMs. In *Findings of the Association for Computational Linguistics: NAACL 2025*. Association for Computational Linguistics, 5541–5558. doi:10.18653/v1/2025.findings-naacl.306
- [56] Yan Zhang and Shibo Li. 2026. ConsistencyGate: Preventing Memory Contamination in LLM Agents via Self-Consistency Admission Control. arXiv:2607.22962 doi:10.48550/arXiv.2607.22962
- [57] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., 46595–46623. doi:10.52202/075280-2020
- [58] Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. 2024. MemoryBank: Enhancing Large Language Models with Long-Term Memory. *Proceedings of the AAAI Conference on Artificial Intelligence* 38, 17 (2024), 19724–19731. doi:10.1609/aaai.v38i17.29946